

## APPENDIX TO CUBIT

### A PROGRAMMING SEMANTICS OF CUBIT

CUBIT complies with the standard specification for database indexes and supports Query, Update, Delete, and Insert operations. It provides the following API.

- `Update(row, val)` retrieves the current value of the specified *row*, updates the value to *val*, and returns *true* if the operation succeeds.
- `Delete(row)` retrieves the current value of the specified *row*, deletes this row, and returns *true* if the deletion succeeds.
- `Insert(val)` appends the value *val* to the tail of the bitmap index, increments global variables like *N\_ROWS*, and returns *true* if the insertion succeeds.
- `Query(val/range)` takes the parameters of both point or range queries, and returns the matching subset of the dataset in the form of either a bitvector or an array of pointers to the underlying tuples.

### B EXISTING UPDATABLE BITMAP INDEXES

**In-place.** The most straightforward approach, denoted *In-place* [7], directly updates the underlying bit-matrix. In order to update the  $k^{th}$  row from value  $v_1$  to  $v_2$ , *In-place* applies the *decode-flip-encode* procedure on both bitvectors of  $v_1$  and  $v_2$ . To delete the  $k^{th}$  row, *In-place* applies the same procedure on the bitvector of the old value and sets its  $k^{th}$  bit from 1 to 0. An insert operation appends bit 1 at the tail of the bitvector of the corresponding value, and then appends bit 0 to the others. *In-place*'s inferior performance comes from the time-consuming decode-flip-encode procedure.

**UCB.** To alleviate the performance issue of *In-place*, UCB [4] introduces an extra bitvector, denoted *existence bitvector* (EB), that indicates whether a given row is valid or not. Initially, all bits in EB are 1s. A delete operation is performed by setting the corresponding bit in EB to 0. An insert operation appends a 1 to the tail of EB, and increments the global variable *N\_ROWS* that indicates the number of rows in UCB. An update operation is transformed to delete-then-append operations, that is, the new value is appended at the tail of the bitvector, and a mapping between the invalidated row ID and the new-appended row ID is kept. By avoiding decoding and then encoding the value bitvectors, UDIs of UCB are supposed to be more efficient than *In-place*. The efficiency of UCB is predicated on EB being highly-compressible. In practice, however, its performance deteriorates sharply as the total number of UDIs performed increases and EB becomes less compressible [2].

**UpBit.** To address the above-discussed issues, the state-of-the-art solution, UpBit [2], maintains a *value bitvector* (VB) and an extra *update bitvector* (UB) for every value in the domain of the indexed attribute. UBs keep track of updates to VBs, that is, UDIs flip bits in UBs that are merged back to VBs in a lazy and batch manner. UBs are highly compressible, resulting in reduced decode-flip-encode overheads. Further details can be found in Section Background of the CUBIT paper and the original UpBit paper [2].

### C PARALLELIZING BITMAP INDEXES

**UpBit.** We parallelize UpBit, the state-of-the-art updatable bitmap index, by using a fine-grained locking mechanism. Specifically, the  $\langle VB, UB \rangle$  pair of every value  $v$  is protected by a reader-writer latch, denoted  $latch_v$ . Global variables like *N\_ROWS* are protected by a global latch  $latch_g$ . Update and delete operations first acquire the  $latch_v$  of all values in shared mode to retrieve the current value of the specified row. Then, they upgrade  $latch_v$  of the corresponding bitvectors to exclusive mode in order to flip the necessary bits. An insert operation acquires  $latch_g$  and the corresponding  $latch_v$  in exclusive mode. Consequently, a query operation acquires  $latch_g$  and the corresponding  $latch_v$  in shared mode.

**UCB.** UCB's UDIs update the only EB, and queries read this EB simultaneously. Therefore, we parallelize UCB by using a global reader-writer latch to synchronize concurrent queries and UDIs. Note that an insert operation holds this latch before updating the global variable *N\_ROWS*.

**In-place.** One way to parallelize *In-place* is to use fine-grained reader-writer latches, the same as in UpBit. However, with this mechanism, an insert operation needs to acquire *cardinality* latches before appending bits to the tail of all the VBs, dramatically reducing the overall throughput. Therefore, we parallelize *In-place* by using a global reader-writer latch, the same as for UCB. Surprisingly, the parallelized *In-place* outperforms UCB for high concurrency (see the evaluation results in our paper).

### D SIZE OF RUBS

In the general case, a RUB has only 0s. As UDIs accumulate, the number of 1s in the RUB increases, so does the size of the RUB (which is stored compactly as a list of positions). We now study the operation sequences that increase the size of a RUB.

**FSM of RUBs.** Conceptually, a RUB is a bit-string with a length equal to the cardinality of the domain, and the  $i^{th}$  bit in this bitvector, denoted  $R_i$ , is associated with the corresponding bit of the  $i^{th}$  VB, denoted  $V_i$ . We study the transition of the RUB by using a Finite-State Machine (FSM), in which each node records the  $\langle R_i, V_i \rangle$  pairs for all possible  $i$ , denoted  $\langle R, V \rangle_i$ . For ease of presentation, except for the initial state (the top-left node indicating that the row is just allocated) and the final state (the top-right node indicating that the row has been deleted), all the  $\langle 0, 0 \rangle$  pairs are removed. Each arrow is labeled with the operation that triggers the transition. For example, an insert operation allocates a new RUB and changes its state from  $\langle 0, 0 \rangle$  to  $\langle 0, 1 \rangle$ , indicating that the corresponding bit of the RUB has been set to 1. An update may change a RUB from  $\langle 0, 0 \rangle, \langle 0, 1 \rangle$  to  $\langle 0, 1 \rangle, \langle 0, 0 \rangle$ , leading to a circular arrow starting from and ending at the same node. That is, there is no transition to a new state because the  $\langle 0, 0 \rangle$  pairs are omitted in the FSM. The complete FSM is shown in Figure 1. We make the following observations.

(1) Except for the bottom-right state, the number of 1s in the RUBs of a row is zero to two with high probability.

(2) The only operation sequence that increases the number of 1s of a RUB is as follows.



updated after a *OpDesc* has been successfully appended to the tail of Delta Log. (2) How shared variables are updated is pre-defined in this *OpDesc* by specifying the old and the new values of each variable. (3) Updating shared variables can be performed by any active threads, such that concurrent threads can help each other complete. (4) Shared variables are updated by only using CAS instructions. (5) No ABA problem can arise. Overall, CUBIT-If guarantees that when a UDI and merge operation owning a *OpDesc* completes, each shared variable (a) has been updated to the specified new value, and (b) has been updated only once.

## F TPC-H

We use the TPC-H benchmark in the evaluation. We now present the details on the experimental setup for running CUBIT over the TPC-H benchmark.

**Dataset.** The DBMS maintains two tables, *ORDERS* and *LINEITEM*. The dates of the tuples in *LINEITEM* span the range of years [1992, 1998], the discounts are distributed in the range [0, 0.1] with increments of 0.01, and the quantities are in the range [1, 50].

**Workloads.** We use the Forecasting Revenue Change Query (Q6) as the scan-intensive query workload. The SQL code for Q6 is listed in Algorithm 1. The value of the first parameter *DATE* is the first of January of a randomly selected year in between [1993, 1997], the parameter *DISCOUNT* is randomly selected within [0.02, 0.09], and the parameter *QUANTITY* is randomly selected within [24, 25].

---

### Algorithm 1: TPC-H Q6.

---

```

1 SELECT sum(l_extensprice × l_discount) as revenue
2 FROM LIMEITEM
3 WHERE l_shipdate >= date'[DATE]'
4   and l_shipdate < date'[DATE]' + interval '1' year
5   and l_discount between [DISCOUNT] ± 0.01
6   and l_quantity < [QUANTITY];

```

---



---

### Algorithm 2: TPC-H RF1.

---

```

7 INSERT a new row into the ORDERS table
8 LOOP random[1, 7] times
9   INSERT a new row into the LINEITEM table
10 END LOOP

```

---



---

### Algorithm 3: TPC-H RF2.

---

```

11 DELETE from ORDERS where o_orderkey = [VALUE]
12 DELETE from LINEITEM where l_orderkey = [VALUE]

```

---

We use New Sales Refresh Function (RF1) and Old Sales Refresh Function (RF2) as the workload of updates. Since our DBMS with CUBIT supports real-time updates to data, it is not necessary for RF1 and RF2 to batch together a number of modifications and then apply them in batch mode. In contrast, each RF1 and RF2 modifies a single tuple in the table *ORDERS* and the corresponding few (1-7) tuples in the table *LINEITEM*, and then updates the three CUBIT instances accordingly, all in one transaction. The SQL code for RF1 and RF2 is listed in Algorithms 2 and 3. According to the TPC-H

specification, the operation distribution of each worker thread is set to 98, 1, and 1 for Q6, RF1, and RF2, respectively.

**CUBIT Instances.** The DBMS creates three CUBIT instances, respectively on the attributes *l\_shipdate*, *l\_discount*, and *l\_quantity*. As a result, each Q6 selects the bitvectors corresponding to 1 of the 7 possible years, 3 of the 11 possible discounts, and 24 or 25 of 50 possible quantities, leading to an average selectivity of  $\frac{1}{7} \times \frac{3}{11} \times \frac{24.5}{50} \approx 2\%$ . We use binning to reduce the number of bitvectors for the attribute Quantity from 25 to 3. Values less than, equal to, and larger than 24 go to one of the three bitvectors, respectively. For each Q6 with CUBIT, the DBMS makes a private copy of one bitvector and then performs bitwise OR/AND operations among 5 (1+3+1) or 6 (1+3+2) bitvectors, to retrieve a list of tuple IDs in *LINEITEM*. The Q6 then fetches these tuples to calculate the final revenue result.

## G TPC-C

We also experiment on TPC-C, by using CUBIT to accelerate a common sub-query that retrieves the ID list of the customers belonging to the specified warehouse and district from the *CUSTOMER* table. We accelerate this sub-query by building indexes on these two attributes using different indexes. The worker threads insert or update an entry after every ten queries to simulate updates. Experimental results show that when the *CUSTOMER* table contains 120K entries, the memory footprint of CUBIT is around 25KB, which is about 320× smaller than Bw-Tree, ART, B<sup>+</sup>-Tree, and Hash indexes. Moreover, CUBIT can steadily return the required ID list in 13 microseconds, 20.5-74.2× faster than the alternatives.

## REFERENCES

- [1] Maya Arbel-Raviv and Trevor Brown. 2018. Harnessing epoch-based reclamation for efficient range queries. *Proceedings of the 23rd ACM SIGPLAN Symposium on PPoPP* (2018).
- [2] Manos Athanassoulis, Zheng Yan, and Stratos Idreos. 2016. UpBit: Scalable In-Memory Updatable Bitmap Indexing. In *Proceedings of the ACM SIGMOD International Conference on Management of Data*. <https://dl.acm.org/citation.cfm?id=2915964>
- [3] Trevor Alexander Brown, William Sigouin, and Dan Alistarh. 2022. PathCAS: an efficient middle ground for concurrent search data structures. *Proceedings of the 27th ACM SIGPLAN Symposium on PPoPP* (2022).
- [4] Guadalupe Canahuat, Michael Gibas, and Hakan Ferhatosmanoglu. 2007. Update Conscious Bitmap Indices. In *Proceedings of the International Conference on Scientific and Statistical Database Management (SSDBM)*. 15–25. <https://doi.org/10.1109/SSDBM.2007.24>
- [5] Maurice Herlihy and Nir Shavit. 2008. *The Art of Multiprocessor Programming*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA.
- [6] Maged M. Michael and Michael L. Scott. 1996. Simple, fast, and practical non-blocking and blocking concurrent queue algorithms. In *PODC '96*.
- [7] Abraham Silberschatz, Henry F. Korth, and S. Sudarshan. 2020. *Database System Concepts, Seventh Edition*.